

Machine Learning

Decision Trees

Dan Goldwasser

dgoldwas@purdue.edu

Goal for Today's class

- Wrap-up KNN and discuss Decision trees
- Intuition about hypothesis space and their expressiveness
- "Real learning": creating a model from data (unlike KNN that just memorized it)
- Overfitting vs. Underfitting in DT

Quick Review: The Hypothesis Space

- What functions should the learning algorithm consider?
 - **Does not** have to search over the same space as $f(x)$
 - Should it be a **more** expressive space? **Less** expressive?
- Does it matter?

Quick Review: Is learning possible?

- How many Boolean functions are there over 4 inputs?
- $2^{16} = 65536$ functions (**why?**)
 - 16 possible outputs.
 - Two possibilities for each output
- Without any data, 2^{16} options
- *Does the data identify the right function?*
- The training data contains 7 examples
 - We still have 2^9 options,

Is learning even possible?

x_1	x_2	x_3	x_4	y
0	0	0	0	?
0	0	0	1	?
0	0	1	0	0 ←
0	0	1	1	1 ←
0	1	0	0	0 ←
0	1	0	1	0 ←
0	1	1	0	0 ←
0	1	1	1	?
1	0	0	0	?
1	0	0	1	1 ←
1	0	1	0	?
1	0	1	1	?
1	1	0	0	0 ←
1	1	0	1	?
1	1	1	0	?
1	1	1	1	?

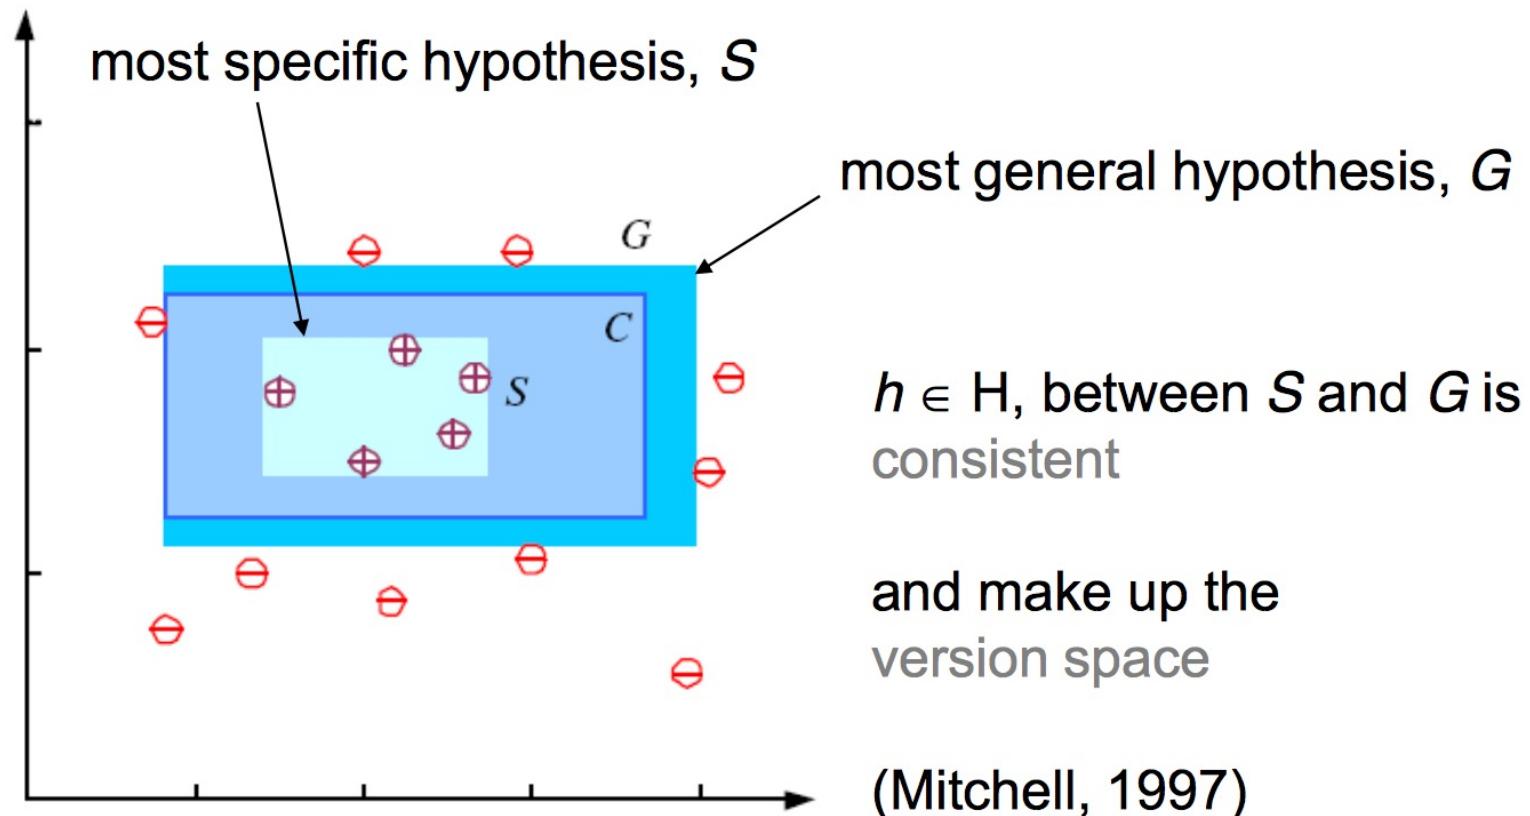
Quick Review: Learning

- Learning is removal of remaining uncertainty
- **Finding a good hypothesis class is essential!**
 - You can start small, and enlarge it until you can find a hypothesis that fits the data (similar argument: “grow” the **feature space**)


Question: *Can there be more than one function that is consistent with the data?* **How do you choose between them?**

A-priori preference towards some functions can be coded in different ways, often via the learning algorithm

Quick Review : hypothesis space and version space




Quick Review: KNN



**Categorical
attributes**

Numerical Attributes:




Distance Measures

$$d(x_1, x_2) = \sqrt{(x_1 - x_2)^2}$$

$$d(x_1, x_2) = 1 - \frac{x_1 \cap x_2}{x_1 \cup x_2}$$

Quick Review: KNN





Quick Review: Let's analyze KNN

- **Time Complexity**
 - *Training is very fast*
 - *Prediction is slower..*
 - *$O(dN)$ for N examples, of dimensionality d*
 - *Increases with number of training examples!*
- **Space Complexity**
 - *KNN needs to maintain all training examples*
- *Let's Compare to the m-of-n rules*

Quick Review: Expressivity

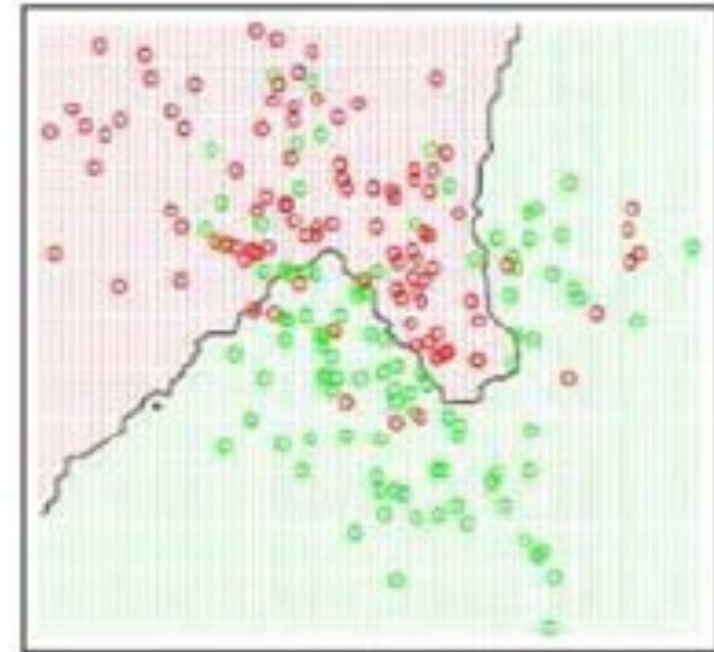

- KNN Can learn very complex decisions models
- We can try to characterize the learned function using its ***decision boundary***
 - *Visualize which elements will be classified as positive/negative*
 - *Decision Boundary is the curve separating the negative and positive regions*



Quick Review: Expressivity

Let's take a closer look at the learned function


→ *High sensitivity to noise!*



Higher k values results in smoother decision boundaries

Figures from Hastie, Tibshirani and Friedman (Elements of Statistical Learning)

Question



Higher k values results in smoother decision boundaries


What will happen if we keep increasing K ?
→ *Extreme scenario: if $k = N$?*

How should we set the value of K?

- **Option 1**: Pick the K that minimizes the **training error**.
 - **Training Error**: number of errors on the training set, *AFTER* we learned a classifier
- **Question**: *What is the training error of 1-NN? 10-NN?*
- **Option 2**: choose k to minimize mistakes on **test error**
 - **Test Error**: number of errors on the test set, after we learned a classifier
 - ***Note***: *It's better to use a validation set, and not the test set. More on that later in the class..*

How should we set the value of K?

How would the test and train error change with K?



In general – using the training error to tune parameters will always result in a more complex hypothesis! **(why?)**

KNN Practical Considerations

- Very simple to understand and implement
- An odd value for k is better (why?)
- How can we find the right value for k ?
 - Using a held out set
- Feature normalization is important! (why?)
 - Different scales for different features

KNN Practical Considerations

- Choosing the right features is important
 - Very sensitive to irrelevant attributes
 - Sensitive to the number of dimensions
- Choosing the right distance metric is important

– Euclidean distance

$$\|\mathbf{x}_1 - \mathbf{x}_2\|_2 = \sqrt{\sum_{i=1}^n (\mathbf{x}_{1,i} - \mathbf{x}_{2,i})^2}$$

– Manhattan distance

$$\|\mathbf{x}_1 - \mathbf{x}_2\|_1 = \sum_{i=1}^n |\mathbf{x}_{1,i} - \mathbf{x}_{2,i}|$$

– L_p -norm

- Euclidean = L_2
- Manhattan = L_1
- Exercise: What is L_∞ ?

$$\|\mathbf{x}_1 - \mathbf{x}_2\|_p = \left(\sum_{i=1}^n |\mathbf{x}_{1,i} - \mathbf{x}_{2,i}|^p \right)^{\frac{1}{p}}$$

Inductive Bias

- Unbiased learning is impossible!
- Inductive bias: **a set of assumptions guiding learning beyond the data**
- We have seen one type of bias: **Language bias**
 - **Pick the right hypothesis space**
- Today, as part of our discussion on Decision Trees we will introduce a second type – **search bias**.
 - **Similar objective:** Encode assumptions about learning, restrict the complexity of the resulting hypothesis

K Nearest Neighbors



- *Is KNN really a learning algorithm?*
 - *Yes, because...*
 - *Well, we get a classifier out of it..*
 - *No, because...*
 - *Well, it just maintains the data..*
- Today we will look at another algorithm that only maintains the data
 - But in a compressed way ("lossy or lossless")

Learning Decision Trees

- KNN only stored the data, Decision trees store a “*compressed*” dataset.
 - Simplified view of DT Learning : *better Compression, with less information loss = better generalization.*
- **DT Learning overview:**
 - Decision Tree Representation
 - Algorithms for learning decision trees
 - Experimental issues
 - Controlling overfitting

Representing Data

- Think about a large table, N attributes, and assume you want to know something about the data represented as entries in the table
 - E.g. own an expensive car or not;
- **Simplest way:** Histogram on the first attribute – own
- **Then,** histogram on first and second (own & gender)




Representing Data

- Think about a large table, N attributes, and assume you want to know something about the data represented as entries in the table
- *But, what if the # of attributes is larger: $N=16$*
 - How large are the 1-d histograms (contingency tables) ? **16 numbers**
 - How large are the 2-d histograms? **$16\text{-choose-2} = 120$ numbers**
 - How many 3-d tables? **560 numbers**
 - **With 100 attributes, the 3-d tables need 161,700 numbers**
- *We need to figure out a way to represent data in a better way, and figure out what are the important attributes to look at first.*
 - we will use an information theoretic approach to represent the data better

Decision Trees



- A hierarchical data structure (tree) that represents data by implementing a divide and conquer strategy
- **Nodes** are tests for feature values
 - There is one branch for every value that the feature can take
- **Leaves** of the tree specify the class labels
- Given a collection of examples, **learn a decision tree that represents it.**
- Use this representation to **classify new examples**
 - The tree can be used for non-parametric classification and regression

Decision Tree Representation



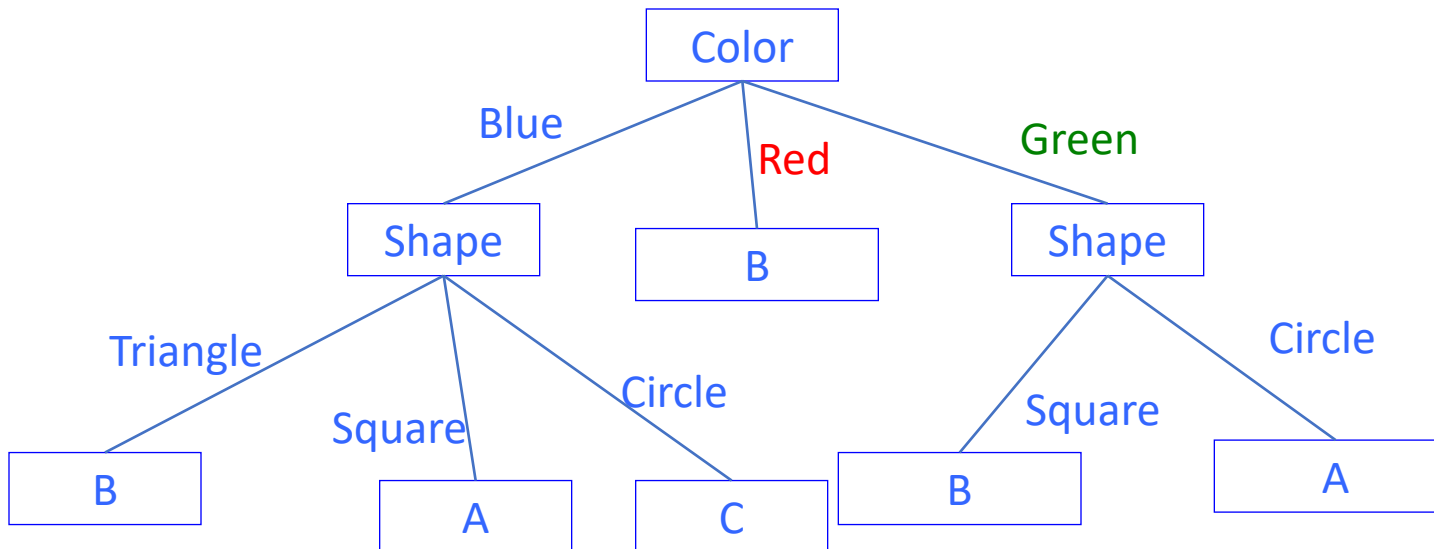
- Three **output** classes: A, B, C
- Two attributes
 - **Color:** Red, Blue, Green
 - **Shape:** Circle, Triangle, Square

Decision Tree Representation



Decision Tree Representation



- **Basic Questions:**
 - *How to use Decision Trees for prediction?*
 - **follow the path from the root**
 - What is the label of a red triangle? Green triangle?
 - **How can we learn decision trees from data?**



Decision Trees


- *Output:*
 - Boolean
 - Multiclass (discrete categories)
 - Real valued (regression tree)
- *Concise representation of the training data*
 - Simply "organizing" the training data in tree form will severely overfit and will be very sensitive to noise (*why?*)
 - Methods for handling noisy data (noise in the label or in the features) and for handling missing attributes
 - **Pruning trees** helps with noise

Expressivity of Decision Trees

- What kind of Boolean functions can DT represent?
 - Any Boolean function! (*why?*)
- A decision tree can be rewritten as a DNF 
 - *Each path from root to leaf can be written as a rule, the tree is a disjunction of these rules.*
 - **Green ^ square → positive**
 - **Blue ^ circle → positive**
 - **Blue ^ square → positive**
- **Question:** *What is the size of the hypothesis space for the shape classification problem?* 

Numeric Attributes

- We have seen instances represented as attribute-value pairs (color=blue, second letter=e, etc.)
 - Values have been categorical
- **What about numeric feature values? (e.g., length)**
 - Discretize them or use thresholds on the numeric values





Expressivity

- What about numeric feature values? (e.g., length)
 - Discretize them or use thresholds on the numeric values
 - *This divides the feature space into axis parallel rectangles*

This is the decision boundary for DT

How would you characterize the complexity of a hypothesis?



Summary: Decision Tree Representation

- *Decision tree can represent any Boolean function*
 - A way to represent lots of data
 - **Compact representation**
 - A natural representation (think 20 questions)
 - “**Explainable**” model!
- *Prediction with a decision tree is easy*
- Clearly, given a dataset, there are many decision trees that can represent it. (**why?**)
 - **How can you learn a good representation from data?**

Example: Will I play tennis today?

- Features

- *Outlook*: {Sun, Overcast, Rain}
- *Temperature*: {Hot, Mild, Cool}
- *Humidity*: {High, Normal, Low}
- *Wind*: {Strong, Weak}

- Labels

- *Binary classification* task: $Y = \{+, -\}$

Basic Decision Tree Learning Algorithm

	O	T	H	W	Play?
1	S	H	H	W	-
2	S	H	H	S	-
3	O	H	H	W	+
4	R	M	H	W	+
5	R	C	N	W	+
6	R	C	N	S	-
7	O	C	N	S	+
8	S	M	H	W	-
9	S	C	N	W	+
10	R	M	N	W	+
11	S	M	N	S	+
12	O	M	H	S	+
13	O	H	N	W	+
14	R	M	H	S	-

Outlook: S(unny),
O(vercast),
R(ainy)

Temperature: H(ot),
M(edium),
C(ool)


Humidity: H(igh),
N(ormal),
L(ow)

Wind: S(trong),
W(eak)

Basic Decision Tree Learning Algorithm

	O	T	H	W	Play?
1	S	H	H	W	-
2	S	H	H	S	-
3	O	H	H	W	+
4	R	M	H	W	+
5	R	C	N	W	+
6	R	C	N	S	-
7	O	C	N	S	+
8	S	M	H	W	-
9	S	C	N	W	+
10	R	M	N	W	+
11	S	M	N	S	+
12	O	M	H	S	+
13	O	H	N	W	+
14	R	M	H	S	-

- Data is processed in Batch (i.e. all the data available)
- Recursively build a DT top down.



Basic Decision Tree Learning Algorithm: ID3

- 1. If all examples are have same label:
 - *Return a single node tree with the label*
- 2. Otherwise
 - Create a **Root node** for tree
 - A = attribute in Attributes that **best** classifies S
 - for each possible value v of attribute A :
 - Add a new **tree branch** corresponding to $A=v$
 - Let S_v be the subset of examples in S with $A=v$
 - if S_v is empty:
 - add **leaf node** with the common value of Label in S
 - Else:
 - below this branch add the **subtree**:
 $ID_3(S_v, \text{Attributes} - \{a\}, \text{Label})$
- 4. Return **Root node**

Input:

S the set of Examples

Label is the target attribute
(the prediction)

Attributes: set of measured
attributes

Why?

Recursive call

Picking the Root Attribute

- The goal is to have the resulting decision tree as small as possible (*Occam's Razor*)
 - But, finding the minimal decision tree consistent with the data is NP- hard
- The recursive algorithm is a greedy heuristic search for a simple tree, but **cannot guarantee optimality**
- The main decision in the algorithm is the selection of the next attribute to split on

Picking the Root Attribute

Consider data with two Boolean Attributes (A,B).


$\langle (A = -0, B = 0), - \rangle$: 50 Examples
 $\langle (A = -0, B = 1), - \rangle$: 50 Examples
 $\langle (A = -1, B = 0), - \rangle$: 0 Examples
 $\langle (A = -1, B = 1), + \rangle$: 100 Examples

What should be the first attribute to split on?

Split on A: purely labeled nodes

Split on B: nodes are not purely labeled

What if we have: “ $\langle (A = 1, B = 0), - \rangle$: 3” examples?



Picking the Root Attribute

Consider data with two Boolean Attributes (A,B).



$\langle (A = -0, B = 0), - \rangle$: 50 Examples

$\langle (A = -0, B = 1), - \rangle$: 50 Examples

$\langle (A = -1, B = 0), - \rangle$: **3 Examples**

$\langle (A = -1, B = 1), + \rangle$: 100 Examples

The trees look structurally similar,
which attributes should you choose?



We need a way to quantify this preference!

Picking the Root Attribute

- The goal is to have the resulting decision tree as small as possible (**Occam's Razor**)
- The main decision in the algorithm is the selection of the next attribute to condition on
- **We want attributes that split the examples to sets that are relatively pure in one label; this way we are closer to a leaf node.**
- The most popular heuristic is based on **information gain**, originated with the ID3 system of Quinlan.

Entropy

Entropy (impurity, disorder) of a set of examples S with respect to binary classification is

$$Entropy(S) = H(S) = -p_+ \log(p_+) - p_- \log(p_-)$$

- *The proportion of positive examples is p_+*
- *The proportion of negative examples is p_-*

In general, for a discrete probability distribution with K possible values, with probabilities $\{p_1, p_2, \dots, p_K\}$ the entropy is given by

$$H(\{p_1, p_2, \dots, p_K\}) = - \sum_{i=1}^K p_i \log(p_i)$$

Entropy

Entropy (impurity, disorder) of a set of examples S with respect to binary classification is


$$Entropy(S) = H(S) = -p_+ \log(p_+) - p_- \log(p_-)$$

- If all examples belong to the same category, entropy = 0
- If $p_+ = p_- = 0.5$, entropy = 1

Entropy

Entropy (impurity, disorder) of a set of examples S with respect to binary classification is

- If all examples belong to the same category, entropy = 0
- If $p_+ = p_- = 0.5$, entropy = 1






Entropy

Entropy (impurity, disorder) of a set of examples S with respect to binary classification is

The uniform distribution has the highest entropy

High Entropy – High level of Uncertainty

Low Entropy – No Uncertainty.



Information Gain

The *information gain* of an attribute A is the *expected reduction in entropy* caused by partitioning on this attribute

$$Gain(S, A) = Entropy(S) - \sum_{v \in \text{Values}(A)} \frac{|S_v|}{|S|} Entropy(S_v)$$

S_v : the subset of examples where the value of attribute A is set to value v

Entropy of partitioning the data is calculated by weighing the entropy of each partition by its size relative to the original set

- partition of low entropy (imbalanced splits) leads to high gain

Go back to check which of the A, B splits is better!

Will I play tennis today?

1	O	T	H	W	Play?
	S	H	H	W	--
2	S	H	H	S	--
3	O	H	H	W	+
4	R	M	H	W	+
5	R	C	N	W	+
6	R	C	N	S	--
7	O	C	N	S	+
8	S	M	H	W	--
9	S	C	N	W	+
10	R	M	N	W	+
11	S	M	N	S	+
12	O	M	H	S	+
13	O	H	N	W	+
14	R	M	H	S	--

Outlook: S(unny),
O(vercast),
R(ainy)

Temperature: H(ot),
M(edium),
C(ool)

Humidity: H(igh),
N(ormal),
L(ow)

Wind: S(trong),
W(eak)

Will I play tennis today?

1	O	T	H	W	Play?
	S	H	H	W	--
2	S	H	H	S	--
3	O	H	H	W	+
4	R	M	H	W	+
5	R	C	N	W	+
6	R	C	N	S	--
7	O	C	N	S	+
8	S	M	H	W	--
9	S	C	N	W	+
10	R	M	N	W	+
11	S	M	N	S	+
12	O	M	H	S	+
13	O	H	N	W	+
14	R	M	H	S	--

Current entropy:

$$p = 9/14$$

$$n = 5/14$$

$$H(Y) =$$

$$-(9/14) \log_2(9/14) - (5/14) \log_2(5/14)$$

$$= 0.94$$

Information Gain: outlook

1	O	T	H	W	Play?
	S	H	H	W	--
2	S	H	H	S	--
3	O	H	H	W	+
4	R	M	H	W	+
5	R	C	N	W	+
6	R	C	N	S	--
7	O	C	N	S	+
8	S	M	H	W	--
9	S	C	N	W	+
10	R	M	N	W	+
11	S	M	N	S	+
12	O	M	H	S	+
13	O	H	N	W	+
14	R	M	H	S	--

Outlook = sunny: 5 of 14 examples

$$p = 2/5 \quad n = 3/5 \quad H_S = 0.971$$

Outlook = overcast: 4 of 14 examples

$$p = 4/4 \quad n = 0 \quad H_o = 0$$

Outlook = rainy: 5 of 14 examples

$$p = 3/5 \quad n = 2/5 \quad H_R = 0.971$$

Expected entropy:

$$(5/14) \times 0.971 + (4/14) \times 0 \\ + (5/14) \times 0.971 = 0.694$$

Information gain:

$$0.940 - 0.694 = 0.246$$

Information Gain: outlook

1	O	T	H	W	Play?
	S	H	H	W	--
2	S	H	H	S	--
3	O	H	H	W	+
4	R	M	H	W	+
5	R	C	N	W	+
6	R	C	N	S	--
7	O	C	N	S	+
8	S	M	H	W	--
9	S	C	N	W	+
10	R	M	N	W	+
11	S	M	N	S	+
12	O	M	H	S	+
13	O	H	N	W	+
14	R	M	H	S	--

Humidity = high:

$$p = 3/7 \quad n = 4/7 \quad H_h = 0.985$$

Humidity = Normal:

$$p = 6/7 \quad n = 1/7 \quad H_o = 0.592$$

Expected entropy:

$$(7/14) \times 0.985 + (7/14) \times 0.592 = \mathbf{0.7885}$$

Information gain:

$$0.940 - 0.7885 = \mathbf{0.1515}$$

Which feature to split on?

1	O	T	H	W	Play?
	S	H	H	W	--
2	S	H	H	S	--
3	O	H	H	W	+
4	R	M	H	W	+
5	R	C	N	W	+
6	R	C	N	S	--
7	O	C	N	S	+
8	S	M	H	W	--
9	S	C	N	W	+
10	R	M	N	W	+
11	S	M	N	S	+
12	O	M	H	S	+
13	O	H	N	W	+
14	R	M	H	S	--

Information gain:

Outlook: 0.246


Humidity: 0.151

Wind: 0.048

Temperature: 0.029

→ *Split on Outlook*

An Illustrative Example




$\text{Gain}(S, \text{Humidity}) = 0.151$

$\text{Gain}(S, \text{Wind}) = 0.048$

$\text{Gain}(S, \text{Temperature}) = 0.029$


$\text{Gain}(S, \text{Outlook}) = 0.246$

An Illustrative Example



	O	T	H	W	Play?
1	S	H	H	W	--
2	S	H	H	S	--
3	O	H	H	W	+
4	R	M	H	W	+
5	R	C	N	W	+
6	R	C	N	S	--
7	O	C	N	S	+
8	S	M	H	W	--
9	S	C	N	W	+
10	R	M	N	W	+
11	S	M	N	S	+
12	O	M	H	S	+
13	O	H	N	W	+
14	R	M	H	S	--

An Illustrative Example




Continue until:

- Every attribute is included in path, or,
- All examples in the leaf have same label

	O	T	H	W	Play?
1	S	H	H	W	--
2	S	H	H	S	--
3	O	H	H	W	+
4	R	M	H	W	+
5	R	C	N	W	+
6	R	C	N	S	--
7	O	C	N	S	+
8	S	M	H	W	--
9	S	C	N	W	+
10	R	M	N	W	+
11	S	M	N	S	+
12	O	M	H	S	+
13	O	H	N	W	+
14	R	M	H	S	--

An Illustrative Example



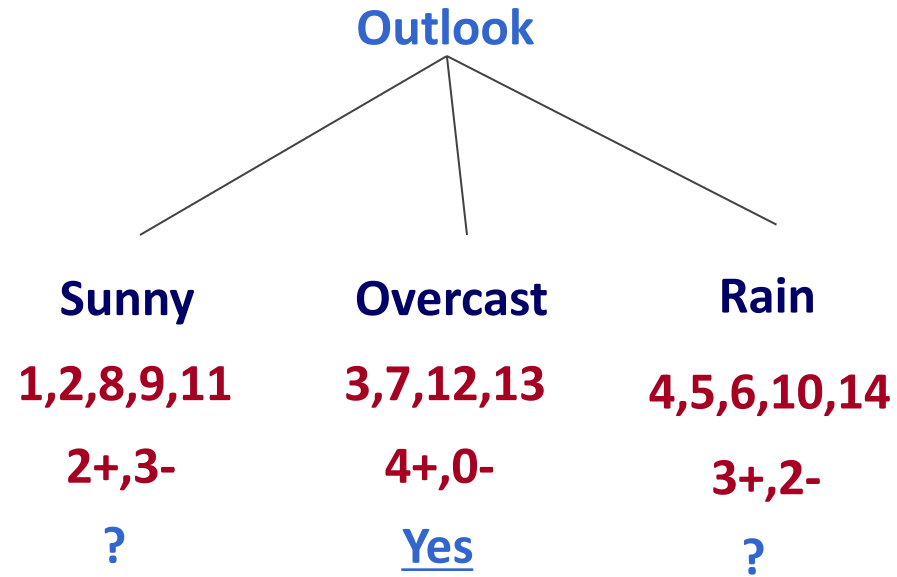
$$\text{Gain}(S_{\text{sunny}}, \text{Humidity}) = .97 - (3/5) \cdot 0 - (2/5) \cdot 0 = .97$$

$$\text{Gain}(S_{\text{sunny}}, \text{Temp}) = .97 - 0 - (2/5) \cdot 1 = .57$$


$$\text{Gain}(S_{\text{sunny}}, \text{wind}) = .97 - (2/5) \cdot 1 - (3/5) \cdot .92 = .02$$

Day	Outlook	Temperature	Humidity	Wind	PlayTennis
1	Sunny	Hot	High	Weak	No
2	Sunny	Hot	High	Strong	No
8	Sunny	Mild	High	Weak	No
9	Sunny	Cool	Normal	Weak	Yes
11	Sunny	Mild	Normal	Strong	Yes


An Illustrative Example



An Illustrative Example



An Illustrative Example



induceDecisionTree(S)

- 1. Does S uniquely define a class?
if all $s \in S$ have the same label y : **return** S ;
- 2. Find the feature with the most information gain:
 $i = \operatorname{argmax}_i \operatorname{Gain}(S, X_i)$
- 3. Add children to S :
for k in $\operatorname{Values}(X_i)$:
 $S_k = \{s \in S \mid x_i = k\}$
 addChild(S, S_k)
 induceDecisionTree(S_k)
return S ;

Variants of Information Gain

- Information gain is defined using entropy to measure the disorder/ impurity of the labels.
- **Other ways to measure disorder**, e.g., *MajorityError*, which computes:
 - *“Suppose the tree was not grown below this node and the majority label were chosen, what would be the error?”*
 - Suppose at some node, there are 15 **+** and 5 **−** examples.
 - What is the *MajorityError*?
 - Answer: $\frac{1}{4}$
- **Similar idea to entropy**

Missing feature values

Suppose an example is missing the value of an attribute. What can we do?


Day	Outlook	Temperature	Humidity	Wind	PlayTennis
1	Sunny	Hot	High	Weak	No
2	Sunny	Hot	High	Strong	No
8	Sunny	Mild	???	Weak	No
9	Sunny	Cool	High	Weak	Yes
11	Sunny	Mild	Normal	Strong	Yes

Missing feature values

- Suppose an example is missing the value of an attribute. What can we do?
- “Complete the example” by
 - Use the most common value of that attribute in data
 - Use the most common value of that attribute in data, with the same class label (at train time)
 - Assign fractional count instead of majority vote on value

Non Boolean Features

- If the features can take multiple values
 - We have seen one edge per value



Non Boolean Features

- If the features can take multiple values
 - We have seen one edge per value (i.e a multiway split)
 - Another option: Make the attributes Boolean by testing for each value

Outlook=Sunny . . [Outlook:Sunny=True,
Outlook:Overcast=False,
Outlook:Rain=False]

Non Boolean Features

What can you do with numeric features?

Use threshold *or* ranges to get Boolean tests

How should you determine the thresholds?


Problem:

You should consider all split points (c) to define node test $X_j > c$

Is there an easier way?

Continuous Attributes


Information gain is minimized when children maintain the same distribution over output labels as the parent node



- ➔ The split should change the proportion of labels in the children
 - ➔ Each child should have a higher proportion one of the label (different)




Continuous Attributes

If a threshold splits two examples of the same label (say, **positive**), and examples on one side of the threshold (say, **left**), have a higher proportion than the parent distribution then:



- *Examples on the other side (right) will have a higher proportion of examples with the other label (negative)*
- *Moving the threshold in this direction (right) until we get to an example with different label, will keep increasing the proportion of that label (positive/negative) in the respective children*

Continuous Attributes



Continuous Attributes

- *The highest information gain split is between examples with different labels.*
 - Simple approach: go over such split points, and find the one with highest information gain
- **Question:**
How many splits should you consider for each continuous attribute?

Bias in decision trees

- Conduct a search of the space of decision trees
 - Can represent all possible functions
- **We prefer short trees!**
 - In DT learning this is implemented as a search bias
 - We bias the search to prefer shorter trees
- Other alternatives?
 - How would you implement language bias on DT?
- Search bias is implemented using greedy heuristics
 - Hill-climbing without backtracking
 - **Overfitting can still be an issue..**

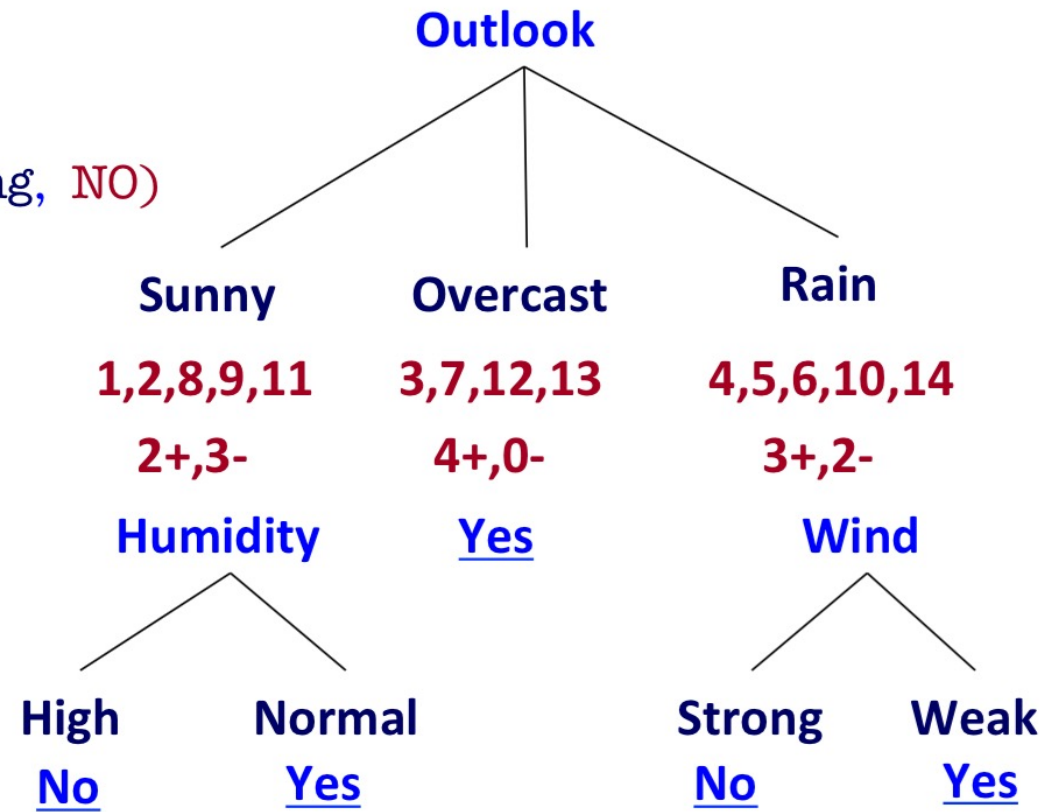
Overfitting

- Learning a tree classifying the training data perfectly may not have the best generalization
 - Algorithm fits tree to noise in training data
 - Sparse data set

A hypothesis h is said to overfit the training data if there is another hypothesis h' , such that h has a smaller error than h' on the training data but h has larger error on the test data.


- Noisy example:

(Outlook = Sunny, Temp = Hot,
Humidity = Normal, Wind = Strong, **NO**)




- Noisy example:

(Outlook = Sunny, Temp = Hot,
Humidity = Normal, Wind = Strong, **NO**)



*may fit noise or other
coincidental regularities*

Decision trees *will* overfit



Avoiding overfitting in decision trees

- **Occam's Razor**

- Favor simpler (in this case, shorter) hypotheses
- Fewer shorter trees, less likely to fit better by coincidence

- **Static:** Fix the depth of the tree

- Only allow trees of size K
 - Tune K using held-out validation set
 - ***Decision stump*** = a decision tree with only one level

- **Dynamic:** optimize while growing the tree

- Grow tree on training data
- Check performance on held-out data after adding a new node

Avoiding overfitting in decision trees

- **Occam's Razor**

- Favor simpler (in this case, shorter) hypotheses
- Fewer shorter trees, less likely to fit better by coincidence

- **Post Pruning:**

- While accuracy on validation set decreases. Bottom up:
 - For each non leaf node:
 - Replace sub-tree under node by a majority vote
 - Test accuracy on validation set

Decision Trees as Features

- When learning over a large number of features, learning decision trees is difficult and the resulting tree may be very large
- **Instead of pruning you can try:**
 - learn small decision trees, with limited depth.
 - Then, learn another function over these trees
- For example, Linear combination of decision stumps

Summary

- **Very popular tool**
 - Prediction is easy (and cheap!)
 - *Expressive and easy to interpret*
 - “debugging” the model is easy!
- **Learning:** greedy heuristic for representing the data
 - ID3 based on information gain
- Prone to **overfitting!**
 - Several ways to deal with it!

Further reading

Machine Learning. Tom Mitchell.

Chapter 3

A Course in Machine Learning. Hal Daumé III.

Chapter 1

(available on line: http://cimpl.info/dl/vo_9/cimpl-vo_9-ch01.pdf)

Questions

- What is *inductive bias*? Why is it important?
- What is the difference between **Language** bias and **search** bias?
- *Which hypothesis space is more expressive?*
 - Boolean functions , Decision trees, linear functions
- How does tree size effect generalization?
- What is the main decision when learning DT?
- How can you deal with noise? Missing attributes? Continuous values?
- *why are decision trees popular?*
 - *When should you use them?*
- *Do you think you can implement a decision tree?*

