

Machine Learning

Multiclass classification and Learning as Optimization

Dan Goldwasser

dgoldwas@purdue.edu

MIDTERM

Midterm!

- Hour long exam.
- 105 possible points
 - 5 bonus points
- **Allowed:** Cheatsheet, calculator, pen
- **Not Allowed:** anything else!

Midterm Question

- Short Question:
 - **True/False:** *“the decision boundary of perceptron and dual perceptron with RBF kernel are similar”*
 - **How would you../what would be the result of doing:**
 - *Reducing the size of a decision tree/changing learning rate.*
- *Short answers (1-2 sentences), consisting of answer (e.g., true/false) and a short explanation*
- ***Avoid guessing.*** *We really just care about the explanation*

Midterm Question

- Calculation Questions:
 - Simulate an algorithm run on some data (small set)
 - Understand the principle behind the algorithm's performance
- Bring a calculator, mostly so you don't waste time.
- If you did the HW you should be fine.
- Make sure answers to questions are consistent with the algorithm "run"

Midterm Questions

- Algorithmic Questions:
 - Adapt the algorithms we have seen to work better in a given setting (little data, high dimensional data, level of noise, etc.)
 - Adapt the algorithms we have seen to new scenarios (come up with new algorithms)
 - We have seen algorithms for learning monotone conjunctions, computing kernels etc. *How can these algorithms be changed?*
 - Analyze algorithms performance

Midterm Questions

- Theoretical Question:
 - Combinatorial questions (counting) and the connection to relevant to ML concepts
 - Show that an algorithm is a mistake bound algorithm
 - Understand the difference between hypothesis classes
- *Make sure to review the definitions and understand them!*


Midterm Topics

- KNN
- Decision Trees
- Learning Boolean Functions
- Winnow
- Perceptron
- Multi-class classification
- Gradient Descent (as far as we'll get today)
- Mistake Bounds analysis, Perceptron convergence
- Performance evaluation, cross-validations, overfitting and underfitting.
- Connections between concepts (how does DT control overfitting?)

- **Explain**. What is one similarity and one difference between..
 - Winnow and Perceptron
- **Pick an option and explain**.
 - Overfitting is more likely in Decision trees or Perceptron.
- **True or False**. A dual perceptron with a linear kernel has the same expressive power as a primal perceptron.
- **True or False**: the size of the hypothesis space (e.g., 3^n for conjunctions, 2^{2^n} for Boolean Functions) is a good indicator for the expressiveness of the space

Example Question


- How would you represent an OR and AND function using a linear threshold function?



- Would a perceptron be able to learn these functions?


Example Question

place points and label them s.t *perceptron will always have zero training error and non-zero leave-one-out validation error*




Example Question

place points and label them s.t perceptron will *always* have zero training error and non-zero leave-one-out validation error




Example Questions

- What will be the result of decision tree learning vs. multi-class perceptron on this dataset?




Example Questions

- What will be the result of decision tree learning vs. multi-class perceptron on this dataset?



Example Questions

- What will be the result of decision tree learning vs. multi-class perceptron on this dataset?



How would it look for Perceptron?
Can you “draw” it on the slide?

Reminder: Loss functions

- To formalize performance let's define a *loss function*:

$$loss(y, \hat{y})$$

- Where \hat{y} is the gold label
- *The loss function measures the error on a single instance*
 - Specific definition depends on the learning task

Regression

$$loss(y, \hat{y}) = (y - \hat{y})^2$$

Binary classification

$$loss(y, \hat{y}) = \begin{cases} 0 & y = \hat{y} \\ 1 & \text{otherwise} \end{cases}$$

Loss minimization

- **Let's consider the square loss.**

- Convex loss function, error surface has a global minimum (~any local minimum is also global).



Do we really want to get to that global minimum point?

- *we care about minimizing the expected loss, while this error surface describes the empirical loss!*


Regularization

- *A form of inductive bias – we prefer simpler functions!*
- A very popular choice of regularization term is to minimize the norm of the weight vector
 - For convenience: $\frac{1}{2}$ squared norm

$$\min_{\mathbf{w}} \sum_n \text{loss}(y_n, \mathbf{w}_n) + \frac{\lambda}{2} \|\mathbf{w}\|^2$$

Loss minimization

- **Let's consider the square loss.**
 - Convex loss function, error surface has a global minimum (~any local minimum is also global).



Gradient Descent for Squared Loss

```
Initialize  $\mathbf{w}^0$  randomly
for  $i = 0 \dots T$ :
     $\Delta \mathbf{w} = (0, \dots, 0)$ 
    for every training item  $d = 1 \dots D$ :
         $f(\mathbf{x}_d) = \mathbf{w}^i \cdot \mathbf{x}_d$ 
        for every component of  $\mathbf{w}$   $j = 0 \dots N$ :
             $\Delta w_j += \alpha(y_d - f(\mathbf{x}_d)) \cdot x_{dj}$ 
     $\mathbf{w}^{i+1} = \mathbf{w}^i + \Delta \mathbf{w}$ 
return  $\mathbf{w}^{i+1}$  when it has converged
```

Stochastic Gradient Descent

Initialize \mathbf{w}^0 randomly

for $m = 0 \dots M$:

$$f(\mathbf{x}_m) = \mathbf{w}^i \cdot \mathbf{x}_m$$

$$\Delta w_j = \alpha(y_d - f(\mathbf{x}_m)) \cdot x_{mj}$$

$$\mathbf{w}^{i+1} = \mathbf{w}^i + \Delta \mathbf{w}$$



return \mathbf{w}^{i+1} when it has converged

Regularization

- Both for regression and classification, for a given error we prefer a *simpler model*
 - *Keep W small: ϵ changes in the input cause $\epsilon * w$ in the output*
- *Some times we are even willing to trade a higher error rate for a simpler model (why?)*
- *Add a regularization term:*
 - *This is a form of **inductive bias***

$$\min_w = \sum_n \text{loss}(y_n, w_n) + \lambda R(w)$$


How different values affect learning?




L1 and L2 norms

- Common choices: L1 and L2 are both convex ($L_{p<1}$ is not convex)
- *Regularized objective*: balance between minimizing the error and the regularization cost

*L2 optimum will be sparse,
ONLY if the data loss term is
minimized at the axis*




*L1 norm contour are **sharp**, will intersect with the contour of data loss term even when the data loss term min point is not at the axis
→ L1 encourages sparsity*




Surrogate Loss functions

- Surrogate loss function: smooth approximation to the 0-1 loss
 - Upper bound to 0-1 loss



Logistic Regression



$$h_w(x) = g(w^T x)$$

$$z = w^T x$$

$$g(z) = \frac{1}{1 + e^{-z}}$$

Sigmoid
(logistic)
function

Known as a sigmoid/logistic function

- Smooth transition between 0-1
- **Can be interpreted as the conditional probability**
- Decision Boundary
 - $y=1: h(x) > 0.5 \rightarrow w^T x \geq 0$
 - $y=0: h(x) < 0.5 \rightarrow w^T x < 0$

- **Learning:** optimize the likelihood of the data
 - **Likelihood:** *probability of our data under current parameters*
 - For easier optimization, we look into the log likelihood (*negative*)


$$Err(w) = - \sum_i y^i \log(e^{g(w, x_i)}) + (1 - y^i) \log(1 - e^{g(w, x_i)}) + \frac{1}{2} \lambda ||w||^2$$

Hinge Loss

- Another popular choice for loss function is the hinge loss

$$L(y, f(x)) = \max(0, 1 - y f(x))$$

- We will discuss in the context of support vector machines (SVM)



It's easy to observe that:

- (1) The hinge loss is an upper bound to the 0-1 loss
- (2) The hinge loss is a good approximation for the 0-1 loss
- (3) BUT ...

It is not differentiable at $y(w^T x) = 1$


Solution: Sub-gradient descent

Reminder: Margin of a classifier

- Distance between a separator (hyperplane) and an example (point)

$$\frac{y(\mathbf{w}^T \mathbf{x})}{\|\mathbf{w}\|}$$

- Margin: the value that minimizes that distance for a given dataset.
- Larger margin can be indicative of better generalization



Hard SVM Intuition

The margin of a classifier: *the distance of the nearest point*

Recall: $\frac{y(\mathbf{w}^T \mathbf{x})}{\|\mathbf{w}\|}$

We want to find the max margin classifier: $\operatorname{argmax}_{\mathbf{w}} [y(\mathbf{w}^T \mathbf{x}) / \|\mathbf{w}\|]$

If we fix $\|\mathbf{w}\| = 1$, we can focus on maximizing the functional margin

$$\mathbf{w}^* = \operatorname{argmax}_{\|\mathbf{w}\|=1} \min_{(x,y) \in S} \overbrace{y(\mathbf{w}^T \mathbf{x})}$$

Or, fix the functional margin $y(\mathbf{w}^T \mathbf{x}) \geq 1$, and focus on minimizing $\|\mathbf{w}\|$

$$\begin{aligned} \mathbf{w}^* &= \operatorname{argmin} \|\mathbf{w}\| \\ \text{s.t. } &y(\mathbf{w}^T \mathbf{x}) \geq 1 \end{aligned}$$

Hard SVM Optimization

- We have shown that the sought-after weight vector \mathbf{w} is the solution of the following optimization problem:

SVM Optimization:

Minimize: $\frac{1}{2} \|\mathbf{w}\|^2$

Subject to: $\forall (x,y) \in S: \quad y \mathbf{w}^T \mathbf{x} \geq 1$

- This is an optimization problem in $(n+1)$ variables, with $|S|=m$ inequality constraints.

Non-Separable Case

Want to relax the constraints:

$$y_i(\mathbf{w} \cdot \mathbf{x}_i + b) \geq 1.$$


Introduce slack variables ξ_i :

$$y_i(\mathbf{w} \cdot \mathbf{x}_i + b) \geq 1 - \xi_i, \quad \text{Where } \xi_i \geq 0, \text{ an error occurs when } \xi_i > 0$$

Thus we can assign an extra cost for errors, as follows:

$$\begin{array}{ll} \text{Minimize} & f(\mathbf{w}, b, \boldsymbol{\xi}) \equiv \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{i=1}^m \xi_i \\ \text{subject to} & y_i(\mathbf{w} \cdot \mathbf{x}_i + b) \geq 1 - \xi_i; \quad \xi_i \geq 0, \quad i = 1, \dots, m \end{array}$$

Visualizing Solution in the non-Separable Case



- | | | |
|-------------------------------|-------------|---------------------|
| 1. Margin support vectors | $\xi_i = 0$ | Correct |
| 2. Non-margin support vectors | $\xi_i < 1$ | Correct (in margin) |
| 3. Non-margin support vectors | $\xi_i > 1$ | Error |

Soft SVM

- Notice that the relaxation of the constraint: $y_i w_i^t x_i \geq 1$ can be done by introducing a slack variable ξ (per example) and requiring:

$$y_i w_i^t x_i \geq 1 - \xi_i \quad ; \quad \xi_i \geq 0$$

- Now, we want to solve:

$$\text{Min } \frac{1}{2} \|w\|^2 + c \sum_i \xi_i \quad \text{subject to } \xi_i \geq 0$$

- ***Which can be written as:***

$$\text{Min } \frac{1}{2} w^T w + c \sum_i \max(0, 1 - y_i w_i^t x_i).$$

Soft SVM (2)

- The hard SVM formulation assumes linearly separability.
 - *A natural relaxation*: maximize the margin while minimizing the # of examples that violate the margin (separability) constraints.
 - *This leads to non-convex problem that is hard to solve.*
 - *Instead*, move to a surrogate loss function that is convex.
- SVM relies on the hinge loss function:
$$\text{Min}_w \frac{1}{2} ||w||^2 + c \sum_i (x,y) \in S \max(0, 1 - y w^t x)$$
- where the parameter c controls the tradeoff between large margin (small $||w||$) and small hinge-loss.

SVM Objective Function

General Form of a learning algorithm:

- **Minimize** empirical loss, and **Regularize** (to avoid over fitting)

$$\text{Min } \frac{1}{2} \|w\|^2 + c \sum \max(0, 1 - y_i w x_i)$$

Regularization term


Empirical loss

Can be replaced by other
regularization functions

Can be replaced by
other **loss functions**

Surrogate Loss functions

- Surrogate loss function: smooth approximation to the 0-1 loss
 - Upper bound to 0-1 loss




Subgradient descent

- asda

Let $f : \mathbb{R}^D \rightarrow \mathbb{R}$ be a convex, not necessarily differentiable, function.

A vector $v \in \mathbb{R}^D$ is called a **subgradient** of f at w_0 , if


$$f(w) \geq f(w_0) + \langle v, w - w_0 \rangle \quad \text{for all } w.$$



Subgradient descent

Let $f : \mathbb{R}^D \rightarrow \mathbb{R}$ be a convex, not necessarily differentiable, function.
A vector $v \in \mathbb{R}^D$ is called a **subgradient** of f at w_0 , if

$$f(w) \geq f(w_0) + \langle v, w - w_0 \rangle \quad \text{for all } w.$$




Subgradient descent

Let $f : \mathbb{R}^D \rightarrow \mathbb{R}$ be a convex, not necessarily differentiable, function.
A vector $v \in \mathbb{R}^D$ is called a **subgradient** of f at w_0 , if

$$f(w) \geq f(w_0) + \langle v, w - w_0 \rangle \quad \text{for all } w.$$

For differentiable f , the gradient $v = \nabla f(w_0)$ is the only subgradient.



Sub-Gradient

Standard 0/1 loss

Penalizes all incorrectly classified examples with the same amount

Hinge loss

Penalizes incorrectly classified examples and correctly classified examples that lie within the margin


Convex,
but not differentiable at $x=1$

Solution: subgradient

The **sub-gradient** of a function c at x_0 is any vector v such that: $\forall x : c(x) - c(x_0) \geq v \cdot (x - x_0)$.

At **differentiable** points this set only contains the gradient at x_0

Intuition: the set of all tangent lines (lines under c , touching c at x_0)



$$\begin{aligned} & \partial_w \max\{0, 1 - y_n(\mathbf{w} \cdot \mathbf{x}_n + b)\} \\ &= \partial_w \begin{cases} 0 & \text{if } y_n(\mathbf{w} \cdot \mathbf{x}_n + b) > 1 \\ y_n(\mathbf{w} \cdot \mathbf{x}_n + b) & \text{otherwise} \end{cases} \\ &= \begin{cases} \partial_w 0 & \text{if } y_n(\mathbf{w} \cdot \mathbf{x}_n + b) > 1 \\ \partial_w y_n(\mathbf{w} \cdot \mathbf{x}_n + b) & \text{otherwise} \end{cases} \\ &= \begin{cases} 0 & \text{if } y_n(\mathbf{w} \cdot \mathbf{x}_n + b) > 1 \\ y_n \mathbf{x}_n & \text{otherwise} \end{cases} \end{aligned}$$